

Filter Representation in Vectorized Query Execution (DAMON'21) — Junior-Friendly Explanation

Paper: “**Filter Representation in Vectorized Query Execution**” (Ngom, Menon, Butrovich, Ma, Lim, Mowry, Pavlo), 2021.

This writeup explains the paper in a way that a junior data engineer can follow. The goal is to understand:

- what “vectorized execution” means
 - what “selection vectors” and “bitmaps” are
 - why one is sometimes faster than the other
 - what the paper’s recommendations are
-

1) The problem (in plain language)

Modern analytical databases often keep data in **memory** (RAM).

That changes the bottleneck:

- Old world: disk was slow → optimize disk reads
- New world: memory is fast → CPU efficiency matters a lot

A popular CPU efficiency trick is **vectorized execution**.

What is vectorized execution?

Instead of processing one row at a time, the engine processes a **batch** (a “vector”) of rows at a time. Typical batch size: ~1–2k tuples.

Why this is faster:

- less function-call / iterator overhead
- better cache usage
- more chances to use **SIMD** (CPU instructions that do the same operation on multiple values at once)

2) Where “filters” come in

During query execution, operators keep narrowing the data:

- scan a table
- apply WHERE conditions
- compute expressions
- joins, groups, etc.

After each step, the system must track:

which rows in the current batch are still “alive” (valid)

The paper calls that tracking structure a **filter representation**.

3) Two filter representations

A) Selection Vector (SV)

A **selection vector** is basically a list of indexes (tuple IDs) that survived.

Example:

- Batch has 8 rows: positions 0 .. 7
- Rows that survived: 0 , 3 , 4 , 7
- Selection vector: [0 , 3 , 4 , 7]

Mental model:

- It's like a guest list: only these people are allowed in.

Cost shape:

- Great when few rows survive (small list).
- But it introduces “indirection”: you read `idx = SV[i]` then access `col[idx]`.

B) Bitmap (BM)

A **bitmap** stores 1 bit per row in the batch.

Example:

- Batch has 8 rows
- Survived rows 0, 3, 4, 7
- Bitmap: 1 0 0 1 1 0 0 1

Mental model:

- It's like a row of light switches; ON means "row is valid."

Cost shape:

- Great when many rows survive (dense selection).
 - Also great when you can process whole vectors with SIMD.
 - But iterating "only the 1s" can be expensive (bit scanning logic).
-

4) Two types of operations: Update vs Map

The paper distinguishes two primitives over filtered vectors:

Update

- Applies a predicate and **changes** which rows are selected.
- Example: WHERE clause filtering rows.

Map

- Computes a new vector but **does not change** which rows are selected.
- Example: projection like SELECT price * 1.2.

Why this distinction matters:

- Some "full compute" strategies are only valid for some combinations.

5) Two compute strategies: Selective vs Full

Given a filter representation, you still have a choice:

Selective compute

Only process rows that are selected.

- If selectivity is low (few rows survive), this saves work.

Full compute

Process all rows in the batch regardless of selection.

- Sounds wasteful, right?
- But it can be faster if the operation is SIMD-friendly and the iteration is cheaper.

This is one of the paper's key messages:

Sometimes doing “more work” is faster because the work is simpler and vectorizes well.

6) Why the “best” choice depends on selectivity

Selectivity = fraction of rows that are selected.

- 0.05 selectivity means 5% of rows survive
- 0.80 selectivity means 80% of rows survive

Think of it like a school bus:

- If only 3 kids are riding, you might drive directly to those 3 houses (selective).
- If 90 kids are riding, you just drive the full route (full), because skipping streets is harder than it's worth.

7) The paper's main finding (the “rule of thumb”)

From their microbenchmarks and analysis:

- **Bitmaps tend to perform better** when the operation can be efficiently **SIMD vectorized**.
- **Selection vectors tend to perform better** for other cases because iteration logic is cheaper.

More detailed:

- SVs are great when selectivity is low (few rows survive) and you want to iterate only survivors cheaply.
 - BMs are great when you can use full scans + SIMD, especially at higher selectivities.
-

8) What the paper actually recommends (practical engineering)

The authors argue that a good vectorized engine should:

1) Support **both** representations (SV and BM) 2) Pick representation/strategy dynamically based on:

- operation cost
- iteration cost
- SIMD friendliness
- selectivity

They show that “one strategy always” is not optimal.

9) Why this matters to you (junior data engineer angle)

Even if you never implement a DBMS, the ideas map to real engineering:

- **Dense vs sparse data structures**
- **Branchy code vs vectorized code**
- **Doing less work vs doing simpler work**

This shows up in:

- analytics engines

- columnar formats
- SIMD-friendly transformations
- vectorized UDF execution

If you work on performance-sensitive pipelines, you'll keep seeing this pattern:

The fastest plan is the one that matches the hardware.

10) Quick glossary

- **Tuple**: a row.
 - **Vector / batch**: a small group of rows processed together.
 - **Selectivity**: fraction of rows that pass a filter.
 - **SIMD**: CPU instructions that operate on multiple values per instruction.
 - **Selection vector**: list of surviving row indexes.
 - **Bitmap**: bitmask of surviving rows.
-

11) Takeaways (memory-friendly)

If you only remember 3 things:

1) Vectorized engines process batches to reduce overhead and use caches better. 2) After each operator, you must track which rows are still valid. 3) SV vs BM is a tradeoff; the best depends on selectivity + SIMD ability.