

# Where Does a Multi■Page Form “Live” Before You Submit It?

If you've ever filled a long form — job application, insurance claim, visa request, checkout flow — you've probably felt this fear:

“If my browser crashes right now... I'm going to become a minimalist and move to the mountains.”

Multi■page forms (often called **wizard forms**) need a place to store your answers **before** they become “official” in the database.

In product/engineering language, that place is usually called **draft storage**.

This article explains draft storage in a way that works for:

- non■technical readers (plain language + analogies)
  - technical readers (correct terms, common architectures, tradeoffs)
- 

## Table of Contents

Table of Contents . . . . .	1
The core idea: draft vs final (email draft analogy) . . . . .	2
Key terms you'll hear (and what they mean) . . . . .	2
What problems draft storage must solve . . . . .	2
Where can the draft live? (the 4 main architectures) . . . . .	3
The most common production pattern: server draft + autosave . . . . .	6
The technical words that matter (and why they matter) . . . . .	6
Security (in plain language) . . . . .	7
Real-world examples . . . . .	8
How to choose (simple decision chart) . . . . .	9

Data engineering angle (why DEs should care) . . . . .	9
Conclusion . . . . .	10

---

## The core idea: draft vs final (email draft analogy)

Think of a multi-page form like writing an email:

- You type and edit → **draft**
- You click “Send” → **final**

The database record is the “sent email.”

Everything before submit is **draft state**. It is real data, but it is not the final record yet.

---

## Key terms you'll hear (and what they mean)

- **Wizard form / multi-step form**: form split into steps/pages.
- **Draft**: a partially completed submission.
- **Autosave**: saving the draft periodically or on step change.
- **Session**: a temporary “conversation” between browser and server (often via cookie).
- **Session store**: where server sessions live (often Redis).
- **TTL (Time To Live)**: automatic expiration time for draft/session data.
- **Idempotency**: retrying a request won’t create duplicates.
- **Optimistic concurrency**: detect “someone else edited this draft” using a version number.
- **PII**: personally identifiable information.

If you can use these words comfortably, you sound like you’ve done this in production (even if you haven’t).

---

## What problems draft storage must solve

A good draft system handles:

- **Next/Back navigation** without losing fields
  - **Refresh / crash / reconnect** without losing progress
  - **Validation** per step and on final submit
  - **Security & privacy** for sensitive fields
  - **Resuming later** (sometimes across devices)
  - **Analytics** (where do users drop off?)
- 

## Where can the draft live? (the 4 main architectures)

### Option 1 — Browser memory (in-memory state)

**Where it lives:** your page's runtime state (React state, Vue state, etc.)

#### Pros

- fastest and simplest

#### Cons

- refresh/close tab → draft is gone

#### Best for

- tiny forms
- low stakes data

Real-world vibe: "If you blink, it forgets."

---

### Option 2 — Browser persistence (localStorage / sessionStorage / IndexedDB)

**Where it lives:** saved on the user's device.

Common options:

- **sessionStorage:** survives refresh, usually dies when tab/browser closes

- **localStorage**: survives browser restart
- **IndexedDB**: better for larger drafts, offline support

### Pros

- refresh-proof, sometimes offline-friendly
- no backend needed for drafts

### Cons

- **security**: if you store sensitive info and you get XSS, it can leak
- device-bound: switching devices won't bring the draft
- users can clear it

### Best for

- “save locally” experiences
- low-to-medium sensitivity data

Analogy: you wrote your draft on a sticky note stuck to *\*this\** laptop.

---

## Option 3 — Server session storage (draft in a session store like Redis)

**Where it lives:** backend, keyed by a session id (often in a cookie).

Typical stack:

- browser has a session cookie
- server stores draft under that session id in **Redis** (or similar)

### Pros

- draft doesn't sit in the browser's localStorage
- central control: TTL, encryption at rest, auditing

### Cons

- session lifecycle complexity

- scaling concerns (session store capacity)
- cross-device resume is tricky unless you tie drafts to a user account

### **Best for**

- authenticated apps
- sensitive data
- when you need strong control and short-lived drafts

Analogy: you're writing your draft on the company's whiteboard (it's safer, but the office locks at night).

---

## **Option 4 — Drafts table (or “staging” records) in the database**

**Where it lives:** a dedicated drafts table (not the final table).

Typical schema:

- `form_drafts(draft_id, user_id, status, current_step, payload_json, created_at, updated_at, expires_at)`

### **Pros**

- very reliable
- easy “resume later” across devices
- supports auditing and compliance

### **Cons**

- you must design cleanup (TTL/expiration)
- partial/incomplete data is normal — don't let downstream treat it as final

### **Best for**

- long, high-stakes forms (gov, insurance, finance)
- regulated workflows
- anything with “Save draft and continue later” as a feature

Analogy: your draft lives in a “draft folder” on the server, not in the “sent mail” folder.

---

## The most common production pattern: server draft + autosave

If I had to bet on what you’ll see in most serious apps:

1) User starts form → server creates a draft and returns `draft_id` 2) Each step (or every N seconds) → client sends partial updates 3) Final submit → server validates everything and writes the final records

### The API usually looks like this

- `POST /drafts` → create draft, returns `draft_id`
- `PATCH /drafts/{draft_id}` → update part of the draft (autosave)
- `GET /drafts/{draft_id}` → resume
- `POST /drafts/{draft_id}/submit` → final validation + write to final tables

Why PATCH? Because you often update “just this step,” not the whole world.

---

## The technical words that matter (and why they matter)

### Idempotency (so retries don’t duplicate)

Networks fail. Mobile networks fail more.

If the client sends “save step 2” twice, you still want **one** draft update.

Common technique:

- include an **Idempotency-Key** header
- or use a `draft_version` / `step_version` field

### Optimistic concurrency (avoid overwriting)

If the user opens the form in two tabs, you don’t want tab A to overwrite tab B silently.

Common technique:

- store a version number in the draft

- require the client to send it back; reject if it's stale

## **TTL + cleanup (your drafts will pile up)**

Drafts are like coffee cups in an office kitchen. If you don't clean them, you'll eventually need a hazmat suit.

Common techniques:

- `expires_at` column + nightly cleanup job
  - Redis TTL for session-based drafts
- 

## **Security (in plain language)**

### **Why localStorage is risky for sensitive data**

If an attacker can run JavaScript on your page (XSS), they can read localStorage.

So don't put:

- passwords
- tokens
- full credit card details
- sensitive PII

If you must store locally:

- minimize what you store
- consider encrypting client-side (but key management becomes its own project)

### **Server-side draft storage helps**

You can:

- encrypt at rest (KMS-managed keys)
- audit access
- enforce ACLs

- expire drafts reliably
- 

## Real-world examples

### Example 1: Government/insurance form

Requirements:

- save for weeks
- resume across devices
- strong auditing

Typical choice:

- **drafts table in DB + autosave**

### Example 2: E-commerce checkout

Requirements:

- fast
- short lived
- handle retries

Typical choice:

- **server session** (Redis) for cart + shipping steps

### Example 3: Job application

Requirements:

- save draft for days
- attachments (CV)

Typical choice:

- **drafts table + object storage for files**

---

## How to choose (simple decision chart)

Ask these questions:

**1) Do users need resume after refresh?**

- No → in-memory is OK
- Yes → localStorage/IndexedDB or server drafts

**2) Do users need resume across devices?**

- Yes → server drafts tied to user

**3) Is the data sensitive?**

- Yes → avoid localStorage; prefer server drafts/session store

**4) How long should drafts live?**

- Minutes/hours → session store with TTL
- Days/weeks → drafts table with `expires_at`

---

## Data engineering angle (why DEs should care)

Draft storage affects:

- **data correctness** (no partial records in final tables)
- **lineage** (what was saved vs submitted)
- **funnel analytics** (drop-off by step)
- **compliance** (retention + deletion)

In other words: draft storage is not just a frontend thing.

## Conclusion

Multi-page forms need a place to put “unfinished work.” That place is draft storage.

If you remember one line:

Keep drafts separate from final records, and treat the draft system like a real data product.

In a future article, I'll go deeper into:

- designing a drafts schema (JSON vs normalized tables)
- autosave performance patterns
- handling file uploads inside multi-step flows
- and the security checklist that prevents very expensive mistakes

If you tell me your use case (public form? logged-in app? sensitive data?), I can recommend the best architecture for your situation. \*\*\* End Patch")Commentary to=functions.apply\_patch ■■ json